

Reinforcement Learning Project Proposal

Joseph (Jack) Bosco
jab2516

Akshara Pramod
ap4613

Date : 11th March, 2025

A Study on Explainability in RL Models

Idea

This project aims to improve explainability in reinforcement learning by incorporating Variational Autoencoders (VAEs) to create meaningful latent representations of the CarRacing-v0 environment. Instead of directly training an RL agent on high-dimensional raw pixel inputs, we propose learning a compressed latent space using a VAE as it would capture essential driving features. The RL model will then learn policies in this reduced representation space, making training more efficient and interpretable. To further enhance explainability, we will apply t-SNE to visualize decision-making in the latent space. We also plan to leverage KL divergence in VAE loss to extract more meaningful and well-formed latent space.

Motivation

For AI models to be safely deployed, especially in environments where human safety or well-being is involved, stakeholders must be able to interpret and trust the decisions made by these models. Deep RL models trained on high-dimensional image data often act as black boxes, making it difficult to interpret why certain decisions are made. This lack of transparency limits trust in RL-based autonomous decision-making, particularly in safety-critical domains like self-driving cars.

Proposal

We aim to follow methods set out in [1] using variational autoencoders (VAE) to learn explicit latent space representations for high-dimensional input space. First we pretrain the autoencoder on a bunch of episodes from an actual human playing the game. The objective is to reconstruct the game frame-by-frame through a low-dimensional bottleneck. After pretraining we append a new neural network to the bottleneck in the architecture which experiences reward from the environment and back-propagates the reward to the encoder. This will be the RL component. For explainability we can use VAE to generate interpretable representations from the state space [2]. More interestingly, we can generate novel images using the geometric average latent space coordinate of a certain feature. This latent space arithmetic is known to reveal model biases and is a useful tool for interpreting how the RL agent responds to certain scenarios.

Dataset: **CarRacing-v2 (OpenAI Gymnasium Box2D)**

The CarRacing-v2 dataset is a continuous high-dimensional control environment from OpenAI Gymnasium's Box2D suite. It involves controlling a car on procedurally generated racetracks, requiring precise navigation and long-term strategy to optimize lap times. The state space consists of 96x96 RGB images, making it a computer vision-based RL problem. The agent controls acceleration, braking, and steering in a continuous action space. The reward is equal to 1000-the time it takes to complete a lap.

Encoder	Decoder
Input: $96 \times 96 \times 3$ RGB image	Input: latent sample $\in \mathbb{R}^k$
Conv. $32 \times 3 \times 3$, stride 2, ReLU, BN	Dense, 256, ReLU
Max pool. 2×2	Dense, $3 \times 3 \times 128$, ReLU reshaped
Conv. $32 \times 3 \times 3$, stride 2, ReLU, BN	Trans. Conv., $128 \times 3 \times 3$, stride 2, ReLU
Conv. $64 \times 3 \times 3$, stride 1, ReLU, BN	Trans. Conv., $64 \times 3 \times 3$, stride 2, ReLU
Avg. pool. 2×2	Trans. Conv., $32 \times 3 \times 3$, stride 2, ReLU
Conv. $128 \times 3 \times 3$, stride 2, ReLU, BN, flatten	Trans. Conv., $32 \times 3 \times 3$, stride 2, ReLU
Dense, 256, ReLU	Trans. Conv., $16 \times 3 \times 3$, stride 2, ReLU
Dense, 2k	Conv. $3 \times 3 \times 3$, stride 1
Output: Diag. Gaussian	Output: Ind. Bernoulli

Table 1: Revised network architecture for a 96x96 RGB input.

References

- [1] Christopher Gebauer and Maren Bennewitz. *The Pitfall of More Powerful Autoencoders in Lidar-Based Navigation*. arXiv:2102.02127 [cs]. Mar. 2021. DOI: [10.48550/arXiv.2102.02127](https://arxiv.org/abs/2102.02127). URL: <http://arxiv.org/abs/2102.02127> (visited on 03/11/2025).
- [2] Tom White. *Sampling Generative Networks*. arXiv:1609.04468 [cs]. Dec. 2016. DOI: [10.48550/arXiv.1609.04468](https://arxiv.org/abs/1609.04468). URL: <http://arxiv.org/abs/1609.04468> (visited on 03/11/2025).